Towards a generic compilation approach for quantum circuits through resynthesis

Arianne Meijer - van de Griend

Abstract—In this paper, we propose a generic quantum circuit resynthesis approach for compilation. We use an intermediate representation consisting of Paulistrings over $\{Z, I\}$ and $\{X, I\}$ called a "mixed ZX-phase polynomial". From this universal representation, we generate a completely new circuit such that all multi-qubit gates (CNOTs) are satisfying a given quantum architecture. Moreover, we attempt to minimize the amount of generated gates.

We show that for large circuits, containing ≥ 100 **Paulistrings, our algorithm generates significantly less CNOTs than previous methods, as well as both the TKET compiler and Qiskit transpiler.**

> Check paper for: use of ", Paritysynth vs paritysynth, runtime vs compilation time, rewrite abstract

I. INTRODUCTION

Small scale quantum computing has recently become accessible to the public through cloud APIs like AWS Braket [2], IBM Qiskit [1], among others. With these new computing platforms comes the need for new compilation methods for quantum programs. Although some classical compilation methods can still be used for quantum computers, some compilation problems require a new solution due to the quantum nature of these devices. For example, most devices do not allow operation between arbitrary qubits. This is generally true of classical computers as well, but there we can solve the problem by copying bits. However, copying qubits is not possible due to the no-cloning theorem [28]. Similarly, optimizing quantum programs is a complex task. We somehow need to understand what the program is doing without accidentally executing it. Because simulating arbitrary quantum programs is infeasible on classical computers. In this paper, we propose a new compilation method for quantum circuits that optimizes the quantum circuit, and in doing so, it generates only quantum operations as allowed by a target quantum computer.

We essentially do this by throwing away the original circuit and generating a completely new circuit. This means that we do not need to cut the circuit into small pieces (as in e.g. [9, 18, 12]) and we do not generate any SWAP gates explicitly (as in [13, 26, 19]). Hence, this is a full-stack re-synthesis approach. Moreover, this is the first resynthesis algorithm that compares itself to common compilers like Qiskit [1] and TKET [24].

Since this approach transforms quantum circuits into quantum circuit, it is technically a transpilation algorithm rather than a compilation algorithm. We will use these terms interchangeably.

Note that the global nature of our proposed approach makes it probably unsuitable as an intermediate step in an existing compilation, unlike for example peep-hole-optimization.

Our algorithm is a combination of three previous algorithms. We will describe these algorithms in more detail in the next section, but for the familiar reader, we give a short overview here. We start with representing a given circuit in an intermediate representation (IR). We use the mixed ZX-phase polynomial IR that was used in [10] which is essentially a sequence of Paulistrings either over the alphabet $\{Z, I\}$ or $\{X, I\}$. This is a universal representation. We synthesize the circuit using the current state-of-the-art phase polynomial synthesis algorithm [25] and optimize the generated CNOTs with PermRowCol [18]. Then, we used simulated annealing to simplify the IR by adding a few CNOTs to the circuit. This last step essentially removes common CNOTs from the original Paulistrings.

Unfortunately, due to the heuristic nature of our algorithm, we cannot make claims about the minimality of newly generated circuits. However, by comparing our algorithm to Qiskit, TKET, and it's subcomponents, we know that the generated circuits have significantly less CNOTs than previous methods.

Our approach is very powerful for circuits generated from multi-qubit Pauli-interactions such as those generated by transforming a Hamiltonian to a quantum circuit.

In the remainder of this paper we will describe quantum circuit re-synthesis problem in Section II, then we describe our algorithm in Section III. The experimentational details can be found in Section IV. The performance of the proposed method is compared in Section V. We give an in-depth discussion of the assumptions we made in the design of this algorithm and how one can or cannot deal with them (Section VI). Lastly, we summarize and point to future research directions in Section VII.

II. PROBLEM DESCRIPTION

The problem that is addressed in this paper is commonly known as the "qubit routing problem" [6]. It is the problem that many quantum computers only support multi-qubit interactions between specific pairs of qubits. This means that not every arbitrary quantum program (quantum circuit) can be immediately executed on an arbitrary quantum computer. The quantum circuit will need to be changed first to accommodate these connectivity constraints (transpilation). Due to the "no-cloning theorem" [28], we are unable to reuse the existing classical methods because they require the copying of quantum states.
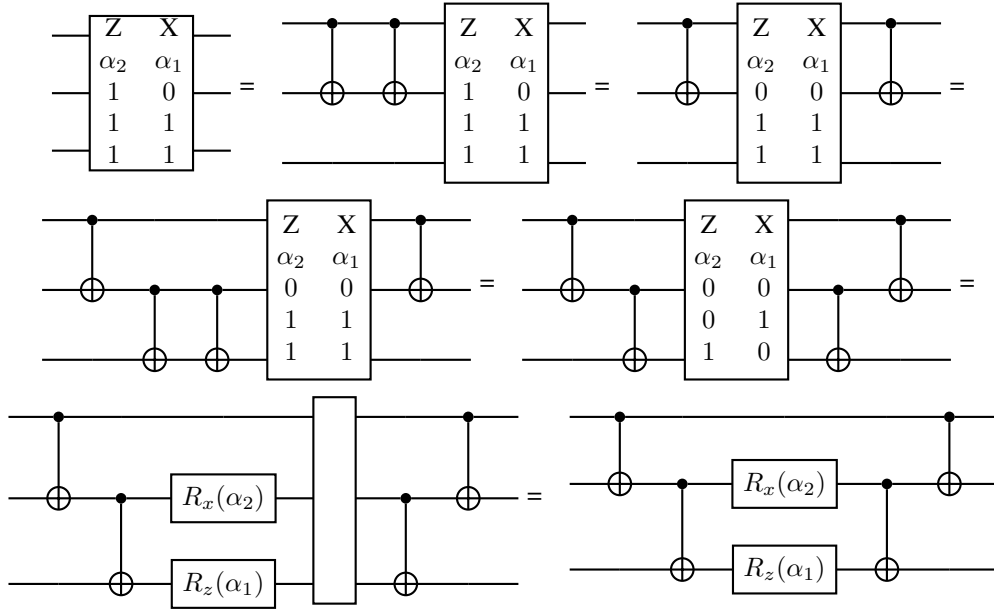
Fig. 1: Example of how the mixed ZX-phase polynomial intermediate representation of $e^{-i\frac{\alpha_2}{2}IXX}e^{-i\frac{\alpha_1}{2}ZZZ}$ can be used to generate a quantum circuit.

As the name suggests, the qubit routing problem is historically solved by routing the qubits: moving the qubits to different registers by applying SWAP operations until the connectivity constraints are satisfied [6]. However, each additional SWAP gate is generally implemented using 3 CNOT gates, which can generate significant overhead for the execution time of the quantum circuit. Since current state-of-the-art quantum computers are noisy, with low decoherence times, and sometimes sparsely connected qubits, the routing overhead can make a circuit impossible to execute.

Thanks to a recent paradigm shift, there is now a class of transpilation methods that do not rely on SWAP gates [21, 12, 9, 17, 18]. Instead, these algorithms rely on finding a scalable representation of the quantum circuit from which a completely new, but equivalent, circuit can be synthesized. During the synthesis process, only multiqubit interactions that are natively allowed are generated, hence the name "architecture-aware resynthesis" [17]. Most of these methods rely on the concept of "Steiner trees" to find how the existing connections can be optimally used. As such, this class of transpilation methods is sometimes called "Steiner-tree-based methods" [18]. An added benefit of resynthesizing the quantum circuit is that it will also be immediately optimized.

A. Mixed ZX-phase polynomial intermediate representation

Our method relies on the intermediate representation named the "mixed ZX-phase polynomial" [10]. We describe it here in detail and how it can be used for resynthesis. Unlike most existing literature, we will not rely on ZX-calculus [27] for this. It is true that mixed ZX-phase polynomials have a very intuitive representation in ZX-calculus, but its use is not yet mainstream enough to assume the reader's familiarity.

Given an arbitrary quantum circuit, we can rewrite that circuit into a sequence of multi-qubit Pauli operations $e^{-i\frac{\alpha}{2}XYZI}$ called "Paulistrings". These operations are usually generated when transforming a Hamiltonian into a quantum circuit, and we assume the reader's familiarity with these operations. Note that a single-qubit gate can also be described as a Paulistring where every letter in the string is I except for one, which corresponds to the axis of rotation for that single-qubit gate.

A circuit consisting only of Paulistrings can then be rewritten into a sequence of Paulistrings either over the alphabet $\{Z, I\}$ (Z phase gadget) or $\{X, I\}$ (X phase gadget). Since Z-phase gadgets mutually commute and X-phase gadgets also, we can group the sequence into sequences of mutually commuting phase gadgets. A sequence of Z-phase gadgets is called a "phase polynomial" [3, 4], hence the name "mixed ZX-phase polynomial".

The advantage of the use of Paulistrings is that they can be efficiently represented. It is fully characterized by the string itself and its angle of rotation α. For the Z- and X-phase gadgets, the string can be represented as a bitstring (given that you know it is a Z- or X-phase gadget). Thus the original quantum circuit can be represented as matrix of $(n+2)\times m$, where n is the number of qubits (+1 for the angle, +1 for whether it is a Z- or X-phase gadget), and m is the number of phase gadgets in the mixed ZX-phase polynomial. Additionally, this representation behaves nicely when CNOTs are acted upon it.

Consider a circuit in which our intermediate representation is a black box. If two consecutive CNOTs were applied at the start of the circuit, the circuit would remain semantically the same. If one were to then move one of the CNOTs through the black box, the intermediate representation changes as follows:

- For **Z-phase gadgets**, the **target** qubit is added to the **control** qubit (modulo 2) in the bitstring.
- For **X-phase gadgets**, the **control** qubit is added to the **target** qubit (modulo 2) in the bitstring.

And the CNOT that was "pushed through" is now behind the black box.

If, at any point during this process, there is a single qubit phase gadget in the *first* sequence of mutually commuting phase gadgets, that phase gadget can be removed from the intermediate representation, and added as a single qubit gate just before the black box.

See Figure 1 for a visual representation of this process. The multi-qubit gate represents the black box with the mixed ZX-phase polynomial representation inside it.

The key observation for this approach to resynthesis is that we need to pick CNOTs that are allowed by the architecture AND optimally reduce the complexity of the intermediate representation in the black box. We want to take advantage of the CNOT being pushed through when it changes the bitstring representations of each phase gadget such that it requires fewer CNOTs later.

In the next section, we will describe how we did this in more detail.

III. THE ALGORITHM

Once the circuit is in the mixed ZX-phase polynomial representation, we need an approximately optimal strategy to synthesize the circuit from it. In this paper, we have combined four previously proposed algorithms [10, 18, 25, 13] into a single procedure. The aim is that these three algorithms work together in such a way that the whole performs better than the separate parts.

First, we use the simulated annealing approach proposed by [10]. The original algorithm adds pairs of CNOTs to the front of the circuit and pushes one of them through the intermediate representation, as explained in the previous section. The simulated annealing algorithm [22] is used to carefully add a few CNOTs to the circuit and simplify the intermediate representation. However, for the final synthesis of the intermediate representation, they use a naive synthesis strategy.

Thus, for this paper, we replace the naive synthesis strategy with a smarter one, namely the algorithm from [25]. We will call this algorithm "Paritysynth" since it tries to synthesize the circuit in an optimal parity ordering for the given connectivity graph. If a subsequence of phase gadgets mutually commutes (because they are of the same type), we can synthesize them in arbitrary order. Thus, we should find the order that is optimal for the CNOT count. This is what Paritysynth does.

Once the intermediate representation in the circuit is empty, we are left with a long sequence of accumulated CNOTs after the black box. These CNOTs can then be optimized using the third algorithm: PermRowCol [18]. The advantage of PermRowCol over other CNOT circuit synthesis procedures [12, 21, 29, 5] is that it is able to reallocate qubits

during synthesis. Like the SWAP-based methods, but without explicitly generating SWAP gates.

This allows us to use the last algorithm: Reverse Traversal [13]. Reverse Traversal was initially used within a SWAP-based compiler, but it can be used in any compilation procedure that reallocates qubits. It uses the fact that quantum computation is reversible. Thus, for every circuit, there exists another circuit that is the reverse of the original circuit. We can run a compilation procedure from an arbitrary initial qubit mapping that reallocates the qubits to different output registers. Then, we can run the compilation procedure on the reverse circuit, using the output mapping of the original circuit as the input for the reverse circuit. The new output mapping can be used as the input mapping for the original circuit, and we can iterate infinitely. If the compilation procedure optimally reallocates the qubits, this should converge to an optimally compiled circuit. Note that Qiskit [1] also uses Reverse Traversal in its compilation procedure.

Since both Reverse Traversal and simulated annealing require multiple iterations of the synthesis procedure, the synthesis procedure needs to be computationally efficient as well. Unfortunately, the Paritysynth algorithm requires the calculation of a Steiner tree for each bit in the bitstring of each phase gadget in the intermediate representation [25]. This can be problematic since finding a minimum Steiner tree is NP-hard [15] and although polynomial approximation exists, the method we used is still cubic with respect to the number of qubits ($O(n^3)$).

To solve this, we have also implemented the phase polynomial synthesis algorithm from [17]. We will call this algorithm "Steiner-Graysynth". Steiner-Graysynth does not rely on the computation of Steiner trees, only the computation of cut vertices, which can be done in $O(n + c)$ using depth-first search, where n is the number of qubits and c the number of edges in the connectivity graph. However, as [25] showed, this method generates more CNOTs than Paritysynth due to the simpler heuristic. We refer the reader to [25] for a detailed discussion.

For sake of brevity, we refer the reader to each of the respective papers for more details.

The novelty of our work is that these algorithms have not yet been used in unison before nor compared against common transpilers like Qiskit and TKET. In the next section, we will see that this combination performs much better than the algorithms on their own.

IV. EXPERIMENTAL SETUP

To evaluate our transpilation procedure, we generated 100 random circuits containing m phase gadgets each acting on at least $\sqrt{n}$ (rounded) qubits, where n is the number of qubits.

The code generating the circuits and implementing our transpilation strategy can be found on Github [1]. Our results are generated using the Jupyter notebook "notebooks/6. Phase Circuit to Quantum Circuit.ipynb". The code was executed on

[1]https://github.com/Aerylia/pauliopt

Fig. 2: Graphs representing the connectivity constraints for different IBM quantum computers. Picture taken from https://www.ibm.com/blogs/research/2019/09/quantum-computation-center/.

a MacBook Air M1 2020 and took approximately 37 hours to complete.

We compare our algorithm against our own implementation of the Steiner-GraySynth algorithm [17] and ParitySynth algorithm [25], the original implementation of the simulated annealing approach [10], the Qiskit transpiler [1], and the TKET compiler [24]. For simulated annealing and reverse traversal, we use 100 iterations, unless they were both used, then they had 10 iterations each to keep total number of iterations the same. Additionally, we used 5 CNOT blocks for the simulated annealer.

As target devices, we used common connectivity constraints of IBM quantum computers. The graphs representing these devices can be found in Figure 2. We chose these devices to give an idea of the scalability of our approach and the effect of different real-world coupling graphs on performance. We encourage readers who are interested in the performance of other graphs to adjust the provided Jupyter notebook for their use cases.

We gathered CNOT counts and runtimes for the following compilation methods:

- **Uncompiled**: This the CNOT count of the original generated circuit where the phase gadgets are naively decomposed. This is the input for the compilation methods and has no associated runtime.
- Compilers:
  - **Qiskit** [1]: Circuit transpiled using Qiskit's "transpile()" method with optimization level 3.
  - **TKET** [24]: Circuit compiled using the default compilation pass with optimization level 2, recreated from documentation [2].
- Baselines:

  - **SG** [17]: Our own implementation of the Steiner-Graysynth algorithm where the CNOT circuit at the end is optimized using PermRowCol [18].
  - **Par** [25]: Our own implementation of the Paritysynth algorithm where the CNOT circuit at the end is optimized using PermRowCol [18].
  - **An** [10]: The original implementation of the simulated annealing approach that uses naive decomposition.
- Proposed methods:
  - **SG+RT**: Uses Reverse Traversal to iteratively improve the SG baseline.
  - **SG+RT→An**: Uses the annealer after SG+RT has found a better initial qubit mapping.
  - **An+RT+SG**: Uses Reverse Traversal to iteratively improve the simulated annealing approach where the annealer uses SG for decomposition.
  - **Par+RT→An**: Same as SG+RT→An but using the Par baseline instead of SG.

Note that we did not run the experiments using the Paritysynth version of 'An+RT+SG' due to runtime issues. We will discuss this in more detail when discussing the runtime results in Table II.

V. RESULTS

We show the average CNOT counts for our randomly generated circuits on the different connectivity graphs in Table I. The column names correspond to the different compilation methods described in section IV.

First, we compare the baselines. The TKET compiler is clearly superior over Qiskit and the annealer. With respect to the other baseline algorithms, TKET is only better for small circuits (1-10 gadgets) and both Steiner-Graysynth and Paritysynth are better than TKET at 100 gadgets. Steiner-Graysynth and Paritysynth seem to perform similarly for small circuits (1-10 gadgets) and the winner seems to be device dependent. However, at 100 gadgets, Paritysynth generates less CNOTs than Steiner-Graysynth. This is in line with results from [25]. The annealer generally generates a few more CNOTs than Qiskit.

When comparing the component baselines with the proposed methods we notice that adding Reverse Traversal to Steiner-Graysynth indeed improves CNOT counts (SG vs. SG+RT). Applying simulated annealer afterward improves the CNOT count even more for circuits with more than 1 phase gadget (SG+RT vs. SG+RT→An). The CNOT count can be lowered even more when the simulated annealing happens as part of the Reverse Traversal iterations (SG+RT→An vs An+SG+RT). However, this comes at a significant increase in the method's runtime, as can be seen in Table II. Nevertheless, the proposed method consistently outperforms Qiskit at this point (Qiskit vs. An+SG+RT), where previously this would only happen for small circuits (Qiskit vs. SG+RT→An).

Moreover, we can improve the proposed annealing methods that use Steiner-Graysynth by replacing it with Paritysynth

Gadgets	Uncompiled	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	5.7	5.46	**3.92**	5.7	*5.96*	5.7	5.44	4.87	4.31	4.81
2	12.1	10.68	**7.93**	11.69	*12.78*	10.8	11.56	10.6	8.8	10.22
5	27.06	23.63	**15.57**	24.22	*25.47*	24.0	23.37	20.46	18.21	20.23
10	50.7	43.09	**27.58**	38.01	38.95	*47.2*	35.21	32.07	29.56	31.87
50	242.44	199.98	121.59	108.25	106.11	*213.48*	104.61	97.45	93.01	**91.33**
100	477.74	392.98	239.91	188.02	187.99	*419.1*	184.17	175.76	169.94	**168.65**

(a) 5-qubit Valencia-like devices

Gadgets	Uncompiled	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	5.44	5.16	**4.6**	5.44	5.44	5.46	5.36	*5.57*	4.98	5.46
2	10.28	9.08	**7.12**	*9.88*	9.39	9.22	9.63	9.07	7.69	9.0
5	24.06	*21.02*	**14.71**	20.41	18.76	19.82	19.16	17.76	15.51	16.3
10	43.98	37.11	25.2	32.32	29.33	*37.66*	29.47	27.32	24.33	**24.23**
50	202.26	*170.64*	112.84	90.0	78.0	155.62	85.5	79.85	75.31	**67.62**
100	416.54	*351.66*	230.88	163.6	139.04	306.76	159.08	149.98	143.47	**125.21**

(b) 5-qubit Yorktown-like devices

Gadgets	Uncompiled	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	19.68	19.35	**14.93**	20.04	*20.42*	20.36	19.8	19.92	18.16	20.11
2	38.8	37.27	**29.09**	38.78	*39.8*	38.6	38.66	37.67	35.2	37.73
5	94.44	88.85	**63.94**	*93.12*	93.02	89.94	93.1	88.05	82.53	86.13
10	183.78	172.09	**123.71**	175.78	174.92	*177.18*	175.64	166.7	158.03	162.0
50	919.24	849.91	622.18	623.2	609.45	*850.76*	598.46	579.1	557.07	**555.07**
100	1833.14	*1690.4*	1223.35	1061.16	1018.56	1688.56	1041.18	1008.54	981.74	**960.73**

(c) 14-qubit Melbourne-like devices

Gadgets	Uncompiled	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	27.1	26.81	**21.33**	*28.9*	27.94	28.1	28.51	28.1	25.55	27.16
2	55.16	53.74	**38.37**	*59.0*	58.05	55.56	58.71	56.25	51.66	54.97
5	137.88	132.82	**92.71**	*146.14*	141.68	134.64	146.09	139.27	129.38	131.21
10	278.36	269.1	**184.88**	*295.36*	273.88	268.84	295.36	280.4	264.57	258.88
50	1366.58	*1311.36*	946.56	1169.5	1095.4	1310.86	1144.14	1100.55	1072.16	1030.03
100	2730.04	*2619.86*	1919.77	1914.16	1780.62	2616.5	1893.95	1842.18	1806.01	**1705.68**

(d) 20-qubit Johannesburg-like devices

Gadgets	Uncompiled	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	27.12	26.8	**20.63**	27.32	*29.04*	27.94	26.81	26.95	24.43	27.34
2	53.46	52.0	**36.1**	53.64	*57.33*	52.94	53.34	51.98	48.72	53.96
5	132.88	126.11	**90.57**	134.7	*141.82*	130.84	134.69	128.12	120.69	131.55
10	270.46	256.31	**180.1**	267.18	*279.74*	261.24	267.18	251.17	240.36	255.55
50	1384.62	1295.42	**938.44**	1090.79	1095.16	*1313.02*	1069.41	1025.26	997.01	1021.67
100	2733.66	2561.52	1904.46	1773.23	1763.2	*2605.04*	1750.49	1693.71	1656.43	**1681.05**

(e) 20-qubit Singapore-like devices

TABLE I: Average CNOT counts over 100 random circuits with different numbers of random phase gadgets and compiled for different devices. Column naming follows the compilation methods as described in section IV.
Numbers in bold are the lowest CNOT counts in the row, and italicized numbers are the highest. Lower counts are better.

(SG+RT→An vs. Par+RT→An). For 100 phase gadget, this even improves the best baseline (TKET vs. Par+RT+→An).

We expect that the CNOT count can still be improved further if the simulated annealing would happen as part of the Reverse Traversal, following the trend for Steiner-Graysynth (SG+RT→An vs An+SG+RT). However, that would increase the runtime to the point that batch experiments like these are infeasible without using a large computational cluster.

Additionally, we note that whenever our best proposed method generates fewer CNOTs than TKET, all other proposed methods do so as well.

The average runtime in seconds of each compilation method is shown in Table II. There is a clear increase in runtime when the circuit has more phase gadgets or the device gets more qubits across all compilation methods. This is expected since these algorithms all run in polynomial time with respect to these factors.

Qiskit is much faster than TKET, but still slower than the other three baseline algorithms. Of the three baselines, Steiner-Graysynth is the fastest, followed by the annealer, and then Paritysynth.

Since Reverse Traversal requires iterative execution, the runtime for Steiner-Graysynth with Reverse Traversal is approximately 100x the runtime for Steiner-Graysynth alone (SG+RT vs. SG). Running the simulated annealer after Reverse Traversal is faster than running Reverse Traversal alone because the algorithm does 10 Reverse Traversal iterations followed by 10 annealer iterations (SG+RT vs SG+RT→An). Running 10 iterations of Reverse Traversal where each iteration runs 10 iterations of the annealer is of course slower (SG+RT→An vs An+SG+RT).

Similarly, the Paritysynth algorithm with Reverse Traversal followed by the annealer is $22-28$x the runtime of Paritysynth alone (Par vs. Par+RT→An). Compared to $27-40$x for

Gadgets	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	0.013	**0.0858**	0.001	*0.0018*	0.0067	0.0915	0.0407	0.317	0.0522
2	0.025	**0.1816**	0.0016	*0.0032*	0.0086	0.1615	0.0634	0.4634	0.0888
5	0.0466	**0.3519**	0.0029	*0.0054*	0.0114	0.2843	0.0964	0.6744	0.1465
10	0.0899	**0.6232**	0.004	0.0085	*0.0149*	0.4031	0.1315	0.9087	0.2256
50	0.4352	2.6337	0.0107	0.0311	*0.0365*	1.2291	0.3201	2.012	**0.7917**
100	0.8999	5.2591	0.0215	0.071	*0.063*	2.4885	0.6062	3.6014	**1.7468**

(a) 5-qubit Valencia-like devices

Gadgets	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	0.0124	**0.096**	0.0009	0.0018	0.0069	0.089	*0.038*	0.2951	0.0499
2	0.0211	**0.1643**	*0.0017*	0.0029	0.0088	0.153	0.0592	0.4389	0.0804
5	*0.0442*	**0.3377**	0.0028	0.0053	0.0123	0.2633	0.0924	0.6606	0.1373
10	0.0753	0.5666	0.0039	0.0079	*0.0157*	0.3879	0.1282	0.8687	**0.207**
50	*0.3587*	2.5011	0.0101	0.0285	0.0374	1.1388	0.3002	1.9016	**0.7144**
100	*0.7246*	5.0354	0.0201	0.0604	0.0618	2.2985	0.5547	3.266	**1.5148**

(b) 5-qubit Yorktown-like devices

Gadgets	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	0.0413	**0.3135**	0.006	*0.0295*	0.0098	0.5552	0.219	1.6737	0.6791
2	0.0804	**0.6688**	0.0116	*0.0532*	0.015	1.0415	0.3836	2.8372	1.2681
5	0.1779	**1.3963**	*0.0224*	0.0971	0.0273	1.9875	0.6941	4.9686	2.3564
10	0.3516	**2.6953**	0.0326	0.1566	*0.0416*	3.2337	1.0666	7.4732	4.0009
50	2.3098	14.5915	0.102	0.7816	*0.1778*	9.8345	2.7383	17.9599	**17.6075**
100	*4.7391*	27.1506	0.1769	1.428	0.2812	18.8626	4.7366	28.4287	**34.4294**

(c) 14-qubit Melbourne-like devices

Gadgets	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	0.0532	**0.4567**	*0.0117*	0.0839	0.0133	1.1052	0.4321	3.2735	1.8964
2	0.1188	**0.9215**	*0.0243*	0.1434	0.0236	2.1109	0.7967	5.8083	3.454
5	0.2735	**2.0804**	*0.0521*	0.3152	0.0463	4.4786	1.5671	11.3152	7.2496
10	0.5497	**4.3389**	*0.0857*	0.537	0.0883	*7.7223*	2.6226	18.5605	12.8408
50	*3.2907*	23.1061	0.2146	2.5235	0.3594	20.9288	5.9286	38.7839	58.1048
100	*6.6193*	46.3972	0.3768	4.9302	0.6347	39.091	10.0422	61.1307	**117.7887**

(d) 20-qubit Johannesburg-like devices

Gadgets	Qiskit [1]	TKET [24]	SG [17]	Par [25]	An [10]	SG+RT	SG+RT→An	An+SG+RT	Par+RT→An
1	0.0546	**0.4602**	0.011	*0.0803*	0.0131	1.0679	0.4066	3.0566	1.8394
2	0.1116	**0.858**	0.0222	*0.137*	0.0217	1.8677	0.6889	5.0574	3.184
5	0.2696	**2.0595**	0.0482	*0.283*	0.0459	4.3044	1.4979	10.5482	6.699
10	0.5531	**4.093**	0.0802	*0.5164*	0.08	7.3885	2.4445	17.2524	12.29
50	3.8567	**21.638**	0.2022	2.4298	*0.3596*	19.9057	5.654	37.3759	56.0521
100	8.1115	**47.5262**	0.3965	5.2028	*0.6972*	38.9048	9.9294	61.696	**116.6315**

(e) 20-qubit Singapore-like devices

TABLE II: Average runtime in seconds over 100 random circuits with different numbers of random phase gadgets and compiled for different devices. Column naming follows the compilation methods as described in section IV.
Bold numbers correspond to the method with the lowest CNOT count in the row and italicized numbers the highest. This is done to more easily see the runtime trade-off for better CNOT counts. Lower runtimes are better.

Steiner-Graysynth (SG vs. SG+RT→An). Then, taking into account that running the annealer as part of the Reverse Traversal is $160 - 315$x the runtume of Steiner-Graysynth, it is reasonable to expect that running a single 20-qubit circuit with 100 phase gadgets can take $5 \cdot 160 = 800s = 13.33$ minutes. This probably takes too long for most use cases but can be worthwhile if the circuit size reduction improves noise or makes it so that the circuit can be executed before decoherence.

More notable is that when our proposed methods generate fewer CNOTs than TKET, some methods also do this faster (SG+RT and SG+RT→An). In fact, SG+RT→An does this in a fraction of the compilation time of TKET and even faster than Qiskit.

In the end, which algorithm is best depends on the use case in which it is applied. Larger circuits benefit more from the synthesis approach than small circuits. How large a circuit needs to be for this trade-off depends on the number of qubits. Circuits with less qubits require less gates for these techniques to outperform TKET. If compilation time is a consideration, TKET and SG+RT→An are good options depending on circuit size. If circuit size is the only consideration, you can increase the number of iterations and go for the Reverse Traversal with Annealing and Paritysynth option that we did not test (it is implemented in our GitHub repository). In that case, one could even run all the compilation methods and take the best result.

VI. DISCUSSION

The proposed approach to the transpilation of quantum circuits relies on a few assumptions. In this section, we list these and discuss why we do not expect them to be a problem.

A. Original circuit quality

Our transpilation approach assumes that the original creator of the quantum circuit to be transpiled is not familiar with quantum circuit optimization and the qubit routing problem. Therefore, we assume that the given circuit is poorly optimized. This is a sensible assumption because the current use cases for quantum computers are the simulation of quantum processes. Thus, the users of quantum computers are physicists and chemists, who might not be familiar with quantum computer science.

As such, we expect that most NISQ programs will be generated from a given Hamiltonian. During this process, the Hamiltonian is transformed into a sequence of Paulistrings, which makes the mixed ZX-phase polynomial a native intermediate representation to the original program. In fact, using Paulistrings as an intermediate representation was already proposed for the Paulihedral compiler [14].

We will mention that if the original quantum circuit is already heavily optimized for the target hardware, we do not expect our transpilation method to still improve the circuit. As we already could see in the results for simple circuits with few phase gadgets. The remainder of the paper will discuss why this could be the case and how one could deal with that.

B. The native multi-qubit interactions of the quantum computer

The only multi-qubit interaction that is generated by our algorithm is the CNOT gate. This makes sense from the point of view that together with arbitrary single qubit gates, this makes a universal gate set. However, the physical implementation of native multi-qubit interactions on current quantum hardware is generally not a CNOT. The devices can execute a CNOT using their native interaction and some single qubit gates, but it is not immediately a CNOT. Some compilation is still necessary!

In general, it takes one of these multiqubit interactions to simulate the CNOT, making it seem that this does not affect the number of multiqubit interactions. However, some devices (e.g. ion-traps) allow multi-qubit interactions between all qubits. This would remove the need to reduce the phase gadgets into smaller interactions.

But let's assume, for the sake of argument, that we have a device with a native 2-qubit interaction. In that case, we can still simulate the CNOT using some single-qubit gates. However, on some devices these interactions will be implemented in an "echoed" form to improve environmental noise [7], doubling all 2-qubit gates. Moreover, in cases where a 2-qubit phase gadget acts on two connected qubits on the topology, the phase gadget might be turned into a native 2-qubit interaction directly (using single-qubit gates) without the need for generating CNOTs.

Transpiling a phase gadget directly to the native 2-qubit gate would require that 2-qubit gate to be tunable with the angle of the phase gadget. Some devices allow this (e.g. some IBM devices [7]), in which case they need to be calibrated for all possible angles rather than just the one angle needed for the CNOT (or two if it is echoed). But calibration of many angles is more difficult. So herein lies a trade-off: implement a 2-qubit phase gadget with 2 "good" interactions (from CNOTs) or 1 "bad" interaction (directly native)?

Our transpilation procedure will change the balance of this trade-off. Due to the resynthesis procedure, the second interaction of the decomposed phase gadget is pushed to the end of the circuit and might be optimized away. On the other hand, if the second interaction isn't optimized away, we might want to use the native interaction. But when the second interaction isn't pushed through the later phase gadget, it will change the entire circuit to the point that it is not obvious which option is better. What is clear in this trade-off is that, because the second interaction might be optimized away, the quality of the tunable 2-qubit interaction needs to be better than twice the fixed-angle interaction.

Additionally, our transpilation approach is not suitable for a look-up-and-replace strategy to introduce the tunable 2-qubit interactions, because it probably will not generate CNOT-(single qubit gate)-CNOT sequences. Thus, in case tunable interactions are allowed, the synthesis approach will need to take that into account from the start.

C. The native single qubit gates

The only single-qubit gates that our transpilation approach generates are X and Z rotations. The assumption here is that these are equally desirable for the hardware, but this is not necessarily true. For example, in superconducting devices, rotations around the Z axis are "virtual". They do not change the qubit itself but the equipment around it (i.e. it resets the clock). This means that Z rotations are essentially noiseless while X rotations require noise manipulations of the qubit. Thus, it is beneficial to generate more Z rotations than X rotations.

Given that X rotations will have to be calibrated and Z rotations do not, it is better to calibrate the X rotations for one or two angles and simulate all arbitrary X rotations $R_X(\alpha)$ as $HR_Z(\alpha)H$, where the H gate is created from the calibrated X rotations and the virtual Z rotation.

This adjustment can be immediately added to our proposed algorithm, but we can also make a slight adjustment to our algorithm that might influence the design of future quantum hardware. During the synthesis procedure, we alternate between decomposing sequences of Z-phase gadgets and X-phase gadgets. What we can do is to generate a H gate on every qubit before the sequence of X-phase gadgets and after the sequence of X-phase gadgets. Due to the connectivity constraints, we do no know which qubits are needed to synthesize the X-phase gadgets, so we need to put the H gate on all qubits (unnecessary H gates can be removed after synthesis). Now, the sequence of X-phase gadgets is turned into a sequence of Z-phase gadgets, and we can continue the algorithm like before, except that when a generated CNOT is pushed through the barrier of H gates, its direction is reversed. In case the sequence of X-phase gadgets is longer than the number of qubits, this reduces the total amount of H gates required.

The reason we believe strategy might influence the design of future quantum hardware is that placing a H gate on all qubits is a very common operation in quantum algorithms [8, 23, 11]. If the transpilation also generates these frequently, it might be sensible to make this operation "hardware accelerated". I.e. make some hardware adjustments so that this specific thing can be done better/faster/easier.

D. Effect of gate fidelity

Another effect that our algorithm does not natively take into account is the difference in gate fidelity. On current quantum computers, not every qubit can be equally well controlled and not every multi-qubit interaction has the same noise. Our results do not take this into account and assume that all gates are equally good. We have discussed some of these aspects when discussing the implementation of the natively supported gates by the quantum devices, but we will discuss the effect of fidelity differences in a broader sense here.

Although we assume for our results that all gates are equally good, our algorithm can be adjusted to take gate quality into account. The most straightforward way to do this is to make the connectivity graph a weighted graph where the weights correspond to the gate fidelities. More generally, we can redefine the measure of "qubit distance" when calculating the minimal Steiner trees. Within our implementation, this would mean redefining the value of the shortest path between every qubit in the topology. Since we approximate the Steiner tree by calculating the minimal spanning tree over the all-pairs shortest paths. This measure of qubit distance can then take more into account than only the 2-qubit gate fidelity, but also the quality of the qubit itself. Similarly, we can use gate fidelity as a tie-breaker when choosing CNOTs or qubits in our algorithm.

In general, we need to find better strategies to take more dynamic sources of noise into account. The algorithm can be adjusted to reduce crosstalk between qubits, improve gate scheduling, or take into account other device-specific sources of error. This will require more collaboration between hardware manufacturers and transpiler designers, as well as the availability of more detailed specifications for the available hardware.

E. Effect of approximate Steiner-trees

Lastly, we discuss the elephant in the room: finding a minimal Steiner tree is an NP-hard problem [15]. The only reason our algorithm is performant is that we use a polynomial approximation. This means that the paths that we find between qubits might not be optimal and we could reduce the CNOT count even further.

It is possible that the types of connectivity graphs on quantum computers are sparse enough so that the approximation is able to find the minimum Steiner tree, but we do not know.

In case the approximate Steiner trees are insufficient, we could still find a minimal Steiner tree using quantum computers. Thanks to adiabatic quantum computing, we should be able to find minimal Steiner trees with quantum annealers [16].

In theory, we could also use gate-based quantum computers to do this, but we run into a chicken-egg problem. Because current quantum annealers have more qubits to their disposal and they do not need a transpiled quantum circuit, we can actually use quantum computers to compile quantum circuits. In fact, this has recently already been done in the ISAAQ compiler [20].

More generally, we might be able to describe the full parity network problem (solved in paritysynth) in a fashion suitable for quantum annealers. There is a lot of interaction between generating which parity on which qubit in which order, which might be very suitable for a quantum annealing formulation.

VII. CONCLUSION AND FUTURE WORK

In this paper, we proposed a full-stack quantum circuit resynthesis approach to transpilation. We compared it to similar state-of-the-art algorithms as well as the Qiskit transpiler and TKET compiler. We showed that for more complicated circuits our proposed approach improves the CNOT count with respect to the baselines. Moreover, it can do so in less compilation time at the cost of more CNOTs.

We argue that in transpilation, we need to take a more holistic approach that takes the semantics of the circuit into account, not just the gates themselves. And we believe that the improved CNOT counts that we have shown support that argument.

However, our algorithm can still be improved. The fact that the simulated annealing is improving the synthesized circuit shows that our heuristics are non-optimal. Moreover, more research is needed on the effect of gate fidelity and approximate Steiner trees. As well as how to make these compiling methods more device-specific, as discussed in Section VI.

Additionally, we need to generalize our algorithm to have Paulistrings as our intermediate representation, just like the Paulihedral compiler [14]. Transforming the Paulistrings into a mixed ZX-phase polynomial can generate many extra phase gadgets which is inefficient.

Lastly, we have shown that there is a trade-off to be made between compilation time and CNOT count. How this trade-off should be made is highly use-case dependent. We believe that the results in our paper and the provided open-source code will make it easier to make such trade-off decisions.

REFERENCES

[1] A-tA-v et al. *Qiskit: An Open-source Framework for Quantum Computing*. 2021. DOI: 10.5281/zenodo.2573505.

[2] Amazon Web Services. *Amazon Braket*. 2023. URL: https://aws.amazon.com/braket/.

[3] Matthew Amy, Parsiad Azimzadeh, and Michele Mosca. "On the controlled-NOT complexity of controlled-NOT–phase circuits". en. In: *Quantum Science and Technology* 4.1 (Sept. 2018), p. 015002. ISSN: 2058-9565. DOI: 10.1088/2058-9565/aad8ca.

[4] Matthew Amy, Dmitri Maslov, and Michele Mosca. "Polynomial-Time T-Depth Optimization of Clifford+T Circuits Via Matroid Partitioning". In: *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* 33.10 (Oct. 2014), pp. 1476–1489. ISSN: 1937-4151. DOI: 10.1109/TCAD.2014.2341953.

[5] Timothée Goubault de Brugière et al. "Quantum CNOT circuits synthesis for NISQ architectures using the decoding problem". In: *Reversible Computation: 12th International Conference, RC 2020, Oslo, Norway, July 9-10, 2020, Proceedings 12*. Springer. 2020, pp. 189–205.

[6] Alexander Cowtan et al. "On the Qubit Routing Problem". en. In: (2019), 32 pages. DOI: 10.4230/LIPICS.TQC.2019.5.

[7] Nathan Earnest, Caroline Tornow, and Daniel J. Egger. "Pulse-Efficient Circuit Transpilation for Quantum Applications on Cross-Resonance-Based Hardware". In: *Physical Review Research* 3.4 (Oct. 2021), p. 043088. ISSN: 2643-1564. DOI: 10.1103/PhysRevResearch.3.043088. URL: https://link.aps.org/doi/10.1103/PhysRevResearch.3.043088 (visited on 10/03/2022).

[8] Edward Farhi, Jeffrey Goldstone, and Sam Gutmann. "A Quantum Approximate Optimization Algorithm". In: (2014). DOI: 10.48550/ARXIV.1411.4028. URL: https://arxiv.org/abs/1411.4028.

[9] Vlad Gheorghiu et al. "Reducing the CNOT Count for Clifford+T Circuits on NISQ Architectures". In: (2020). DOI: 10.48550/ARXIV.2011.12191. URL: https://arxiv.org/abs/2011.12191 (visited on 10/03/2022).

[10] Stefano Gogioso and Richie Yeung. "Annealing Optimisation of Mixed ZX Phase Circuits". In: (2022). DOI: 10.48550/ARXIV.2206.11839. URL: https://arxiv.org/abs/2206.11839 (visited on 10/03/2022).

[11] Lov K. Grover. "A fast quantum mechanical algorithm for database search". en. In: *Proceedings of the twenty-eighth annual ACM symposium on Theory of computing - STOC '96*. Philadelphia, Pennsylvania, United States: ACM Press, 1996, pp. 212–219. ISBN: 978-0-89791-785-8. DOI: 10.1145/237814.237866. URL: http://portal.acm.org/citation.cfm?doid=237814.237866.

[12] Aleks Kissinger and Arianne Meijer van de Griend. "CNOT Circuit Extraction for Topologically-Constrained Quantum Memories". In: *Quantum Information and Computation* 20.7&8 (June 2020), pp. 581–596. ISSN: 15337146, 15337146. DOI: 10.26421/QIC20.7-8-4. URL: http://www.rintonpress.com/journals/doi/QIC20.7-8-4.html (visited on 10/03/2022).

[13] Gushu Li, Yufei Ding, and Yuan Xie. "Tackling the qubit mapping problem for NISQ-era quantum devices". In: *Proceedings of the Twenty-Fourth International Conference on Architectural Support for Programming Languages and Operating Systems*. 2019, pp. 1001–1014.

[14] Gushu Li et al. "Paulihedral: A Generalized Block-Wise Compiler Optimization Framework for Quantum Simulation Kernels". In: *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*. Lausanne Switzerland: ACM, Feb. 2022, pp. 554–569. ISBN: 978-1-4503-9205-1. DOI: 10.1145/3503222.3507715. URL: https://dl.acm.org/doi/10.1145/3503222.3507715 (visited on 10/03/2022).

[15] Ivana Ljubić. "Solving Steiner trees: Recent advances, challenges, and perspectives". en. In: *Networks* 77.2 (Mar. 2021), pp. 177–204. ISSN: 0028-3045, 1097-0037. DOI: 10.1002/net.22005.

[16] Andrew Lucas. "Ising formulations of many NP problems". In: *Frontiers in Physics* 2 (2014). ISSN: 2296-424X. DOI: 10.3389/fphy.2014.00005. URL: http://journal.frontiersin.org/article/10.3389/fphy.2014.00005/abstract.

[17] Adriana Meijer - van de Griend and Ross Duncan. "Architecture-Aware Synthesis of Phase Polynomials for NISQ Devices". In: (2020). DOI: 10.48550/ARXIV.2004.06052. URL: https://arxiv.org/abs/2004.06052 (visited on 10/03/2022).

[18] Adriana Meijer - van de Griend and Sarah Meng Li. "Dynamic Qubit Allocation and Routing for Constrained Topologies by CNOT Circuit Re-Synthesis". In: (2022). DOI: 10.48550/ARXIV.2205.00724. URL: https://arxiv.org/abs/2205.00724 (visited on 10/03/2022).

[19] Prakash Murali et al. "Formal constraint-based compilation for noisy intermediate-scale quantum systems". en. In: *Microprocessors and Microsystems* 66 (Apr. 2019), pp. 102–112. ISSN: 01419331. DOI: 10.1016/j.micpro.2019.02.005.

[20] Soshun Naito et al. "ISAAQ: Ising Machine Assisted Quantum Compiler". In: *arXiv preprint arXiv:2303.02830* (2023).

[21] Beatrice Nash, Vlad Gheorghiu, and Michele Mosca. "Quantum circuit optimizations for NISQ architectures". In: *Quantum Science and Technology* 5.2 (2020), p. 025010.

[22] Martin Pincus. "A Monte Carlo Method for the Approximate Solution of Certain Types of Constrained Optimization Problems". en. In: *Operations Research* 18.6 (Dec. 1970), pp. 1225–1228. ISSN: 0030-364X, 1526-5463. DOI: 10.1287/opre.18.6.1225.

[23] P.W. Shor. "Algorithms for quantum computation: discrete logarithms and factoring". In: *Proceedings 35th Annual Symposium on Foundations of Computer Science*. Santa Fe, NM, USA: IEEE Comput. Soc. Press, 1994, pp. 124–134. ISBN: 978-0-8186-6580-6. DOI: 10.1109/SFCS.1994.365700. URL: http://ieeexplore.ieee.org/document/365700/.

[24] Seyon Sivarajah et al. "t—ket⟩: a retargetable compiler for NISQ devices". In: 6 (Jan. 2021), p. 014003. ISSN: 2058-9565. DOI: 10.1088/2058-9565/ab8e92.

[25] Vivien Vandaele, Simon Martiel, and Timothée Goubault de Brugière. "Phase Polynomials Synthesis Algorithms for NISQ Architectures and Beyond". In: *Quantum Science and Technology* 7.4 (Oct. 2022), p. 045027. ISSN: 2058-9565. DOI: 10.1088/2058-9565/ac5a0e. URL: https://iopscience.iop.org/article/10.1088/2058-9565/ac5a0e (visited on 09/29/2022).

[26] Davide Venturelli et al. "Compiling quantum circuits to realistic hardware architectures using temporal planners". In: *Quantum Science and Technology* 3.2 (Apr. 2018), p. 025004. ISSN: 2058-9565. DOI: 10.1088/2058-9565/aaa331.

[27] John van de Wetering. "ZX-calculus for the working quantum computer scientist". In: *arXiv preprint arXiv:2012.13966* (2020).

[28] William K Wootters and Wojciech H Zurek. "A single quantum cannot be cloned". In: *Nature* 299 (1982), pp. 802–803.

[29] Bujiao Wu et al. "Optimization of CNOT circuits under topological constraints". In: *arXiv preprint arXiv:1910.14478* (2019).